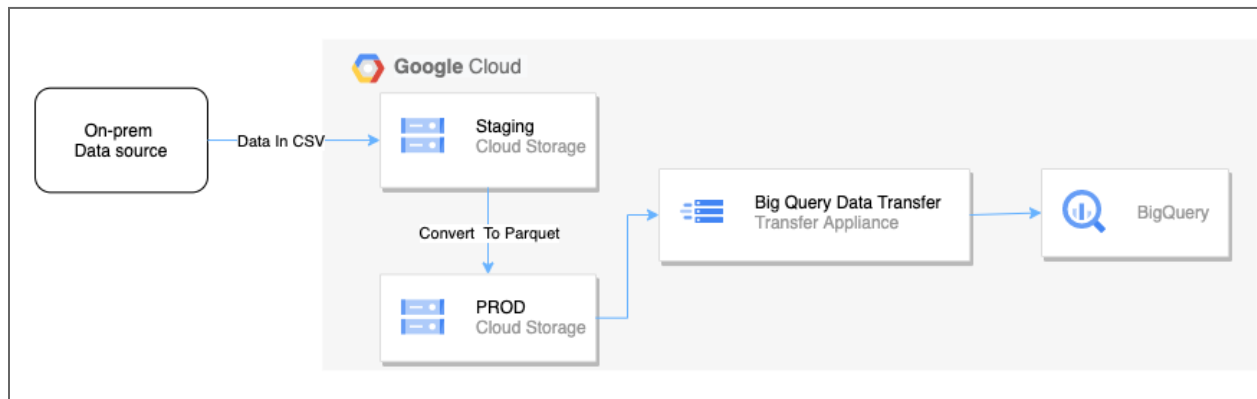


Using BigQuery Data Transfer Service

Here is a working solution which explains a working demo for data load from Google Cloud Storage to BigQuery Data Transfer Service.



Here are the following Steps that load the data to Big query from GCS using Big Query.

Prerequisites

You must have billing enabled in your GCP project. If not sure, follow these [instructions](#).

APIs to be enabled

BigQuery API
BigQuery Data Transfer API

IAM Roles and Permissions


The only predefined IAM roles with all the BigQuery and Cloud Storage permissions are **BigQuery Admin** (roles/bigquery.admin) and **Storage Admin** (roles/storage.admin)

Follow these [instructions](#) to grant the IAM roles.

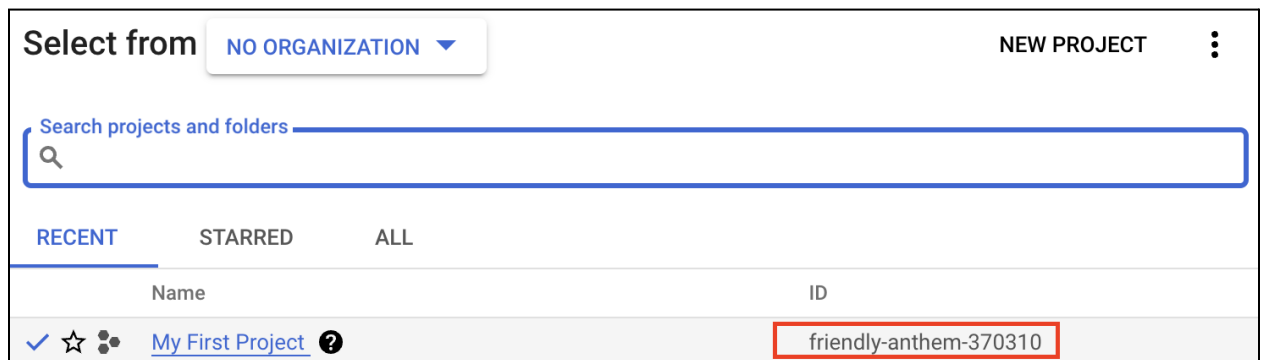
Steps to setup development environment on Cloud Shell Editor

Prerequisites: Gather project ID


On the GCP Console top bar, click on project name (besides Google Cloud logo) e.g. My First Project is referred to in the image below.



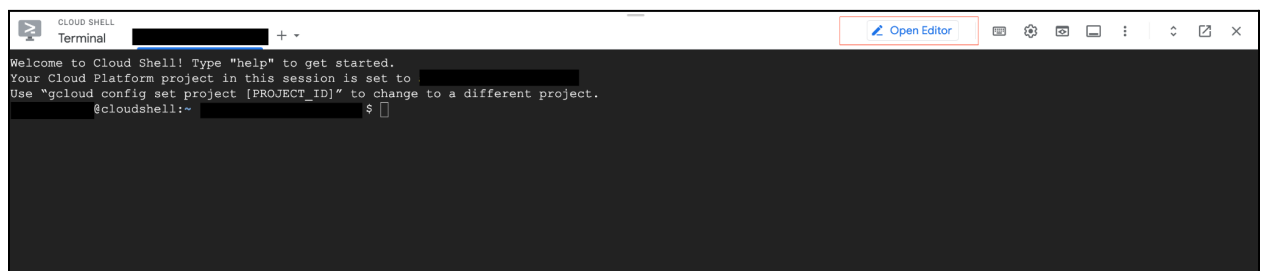
Pop up will appear which contains Project **Name** and Project **ID**. Save the project ID for further usage.



You can use Cloud Shell Editor to run lab's commands. On the GCP Console, on top right click on **Activate Cloud Shell**.



Cloud Shell terminal will open at the bottom of the screen. Click on the **Open Editor** button to access the Cloud Shell Editor.

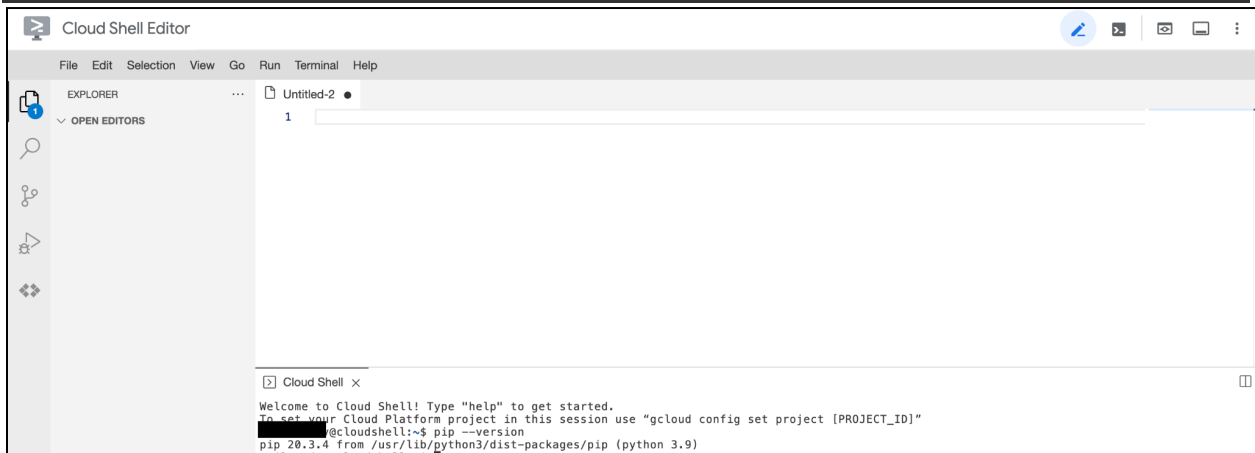


Cloud Shell Editor will be opened refer to below screenshot. Click on **Terminal** to open terminal to run various operations.



Run the following commands on the cloud shell editor terminal to check if you have PIP already installed.

```
pip --version
```



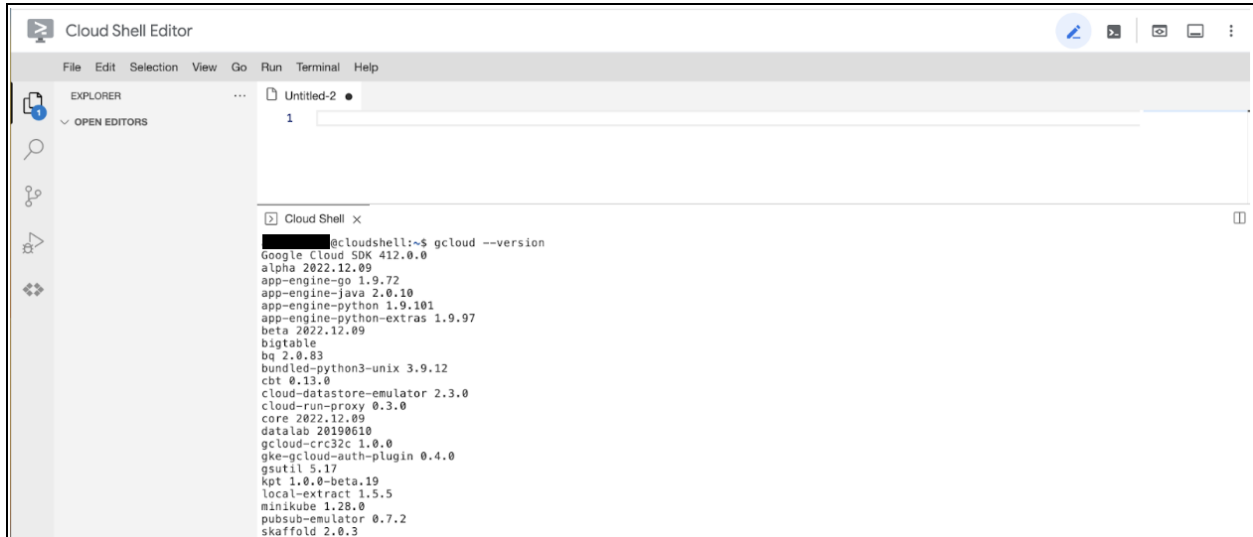
Note: If it is not installed, follow these instructions to [Install](#) PIP according to the operating system you are using.

To run the Python scripts to interact with GCP resources from cloud shell editor, you will need to set up a development environment using the following steps.

- I. Run the following commands on cloud shell editor to check if you have the Google Cloud SDK installed.

```
gcloud --version
```

If you get an output similar to the screenshot below, your SDK installation is done.



Note: If it is not installed, follow these instructions to [Install](#) and [Initialize](#) Google Cloud CLI according to the operating system you are using.

- II. Run the following command on the cloud shell editor terminal to set the current working GCP project

```
gcloud config set project PROJECT_ID
```

Make sure you replace PROJECT_ID with the ID of your GCP Project.

Download data for testing

For this workshop we will be using a public Kaggle Dataset on Netflix Movies and TV Shows.

File Size: 3.4 MB

Number of Rows: 8807

Number of Columns: 12

Link to download the dataset: <https://www.kaggle.com/datasets/shivamb/netflix-shows>

Note: You may have to login or sign up to Kaggle to download the data.

Save it to your computer, extract the zip file and rename it as `netflix_titles.csv`

Create GCS buckets and upload the data

Create two Google Cloud Storage Buckets

- synopsys-staging
- synopsys-prod


To create the buckets:

1. In the Google Cloud console, go to the Cloud Storage **Buckets** page.
[Go to Buckets](#)
2. Click **Create bucket**.
3. On the **Create a bucket** page, enter your bucket information. To go to the next step, click **Continue**.
 - For **Name your bucket**, you must enter values that are globally unique. Since 'synopsys-staging' and 'synopsys-prod' are already taken, add a suffix to make it unique. Eg 'synopsys-staging-323' or 'synopsys-prod-982'
 - For **Choose where to store your data**, select **Location type** as **Region** and select the region that is closest to you from the dropdown list.

Note: You must replace 'synopsys-staging' and 'synopsys-prod' with the bucket names you created for the rest of the tutorial.

Leave the other options as Default and click CREATE.

Navigate to your staging bucket and click the UPLOAD FILES button to upload the 'netflix_titles.csv' file.



This is how your bucket should look once the file is uploaded.

Buckets > synopsys-staging										
UPLOAD FILES UPLOAD FOLDER CREATE FOLDER TRANSFER DATA MANAGE HOLDS DOWNLOAD DELETE										
Filter by name prefix only Filter objects and folders Show deleted data										
<input type="checkbox"/>	Name	Size	Type	Created	Storage class	Last modified	Public access	Version history	Encryption	Retention expiration
<input type="checkbox"/>	netflix_titles.csv	3.2 MB	text/csv	Nov 15, 2...	Standard	Nov 15, 20...	Not public	—	Google-managed key	—

Steps to load CSV data from GCS to BigQuery

The following steps will be used to transfer the CSV file to another GCS Bucket after converting it to Parquet.

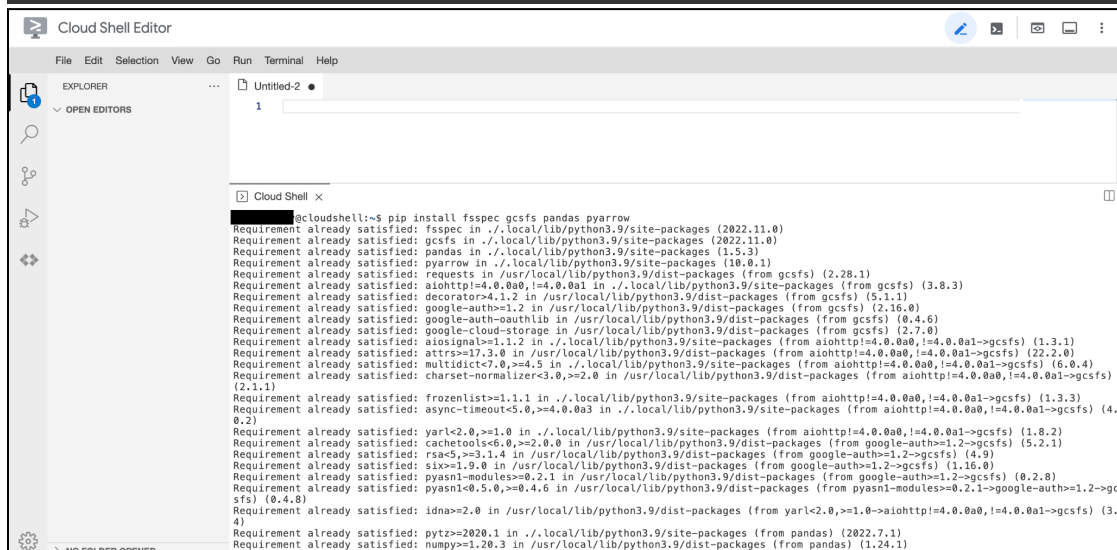
What is Parquet?

Apache Parquet is an open source, column-oriented data file format designed for efficient data storage and retrieval. It provides efficient data compression and encoding schemes with enhanced performance to handle complex data in bulk. We will be converting the large CSV file to Parquet for faster and more efficient processing.

Step 1: CSV to Parquet


- Open cloud shell Editor.
- Run the following command on your cloud shell Editor. terminal to install the prerequisite packages

```
pip install fsspec gcsfs pandas pyarrow
```



The screenshot shows the Cloud Shell Editor interface with a terminal window. The terminal output displays the command `pip install fsspec gcsfs pandas pyarrow` and a list of requirements that are already satisfied, including `fsspec`, `gcsfs`, `pandas`, and `pyarrow`.

- Create a new file in cloud shell editor. Refer below screenshot for reference.



- d. Copy and paste below python code in the newly created file. Save the file with the name 'csv_parquet.py'.

```
import pandas as pd

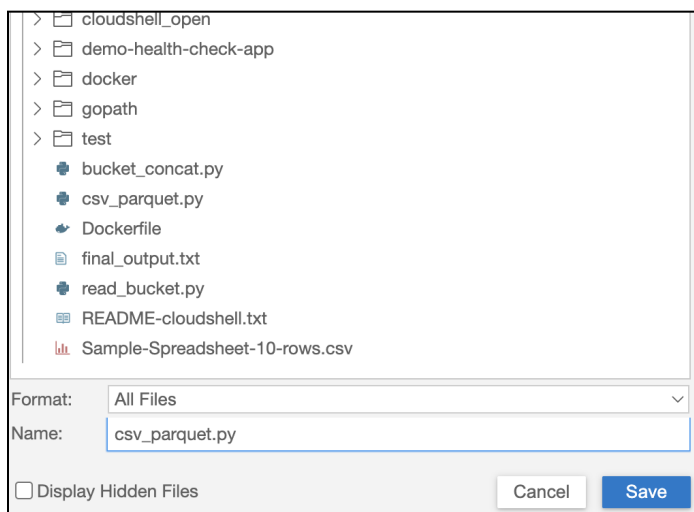
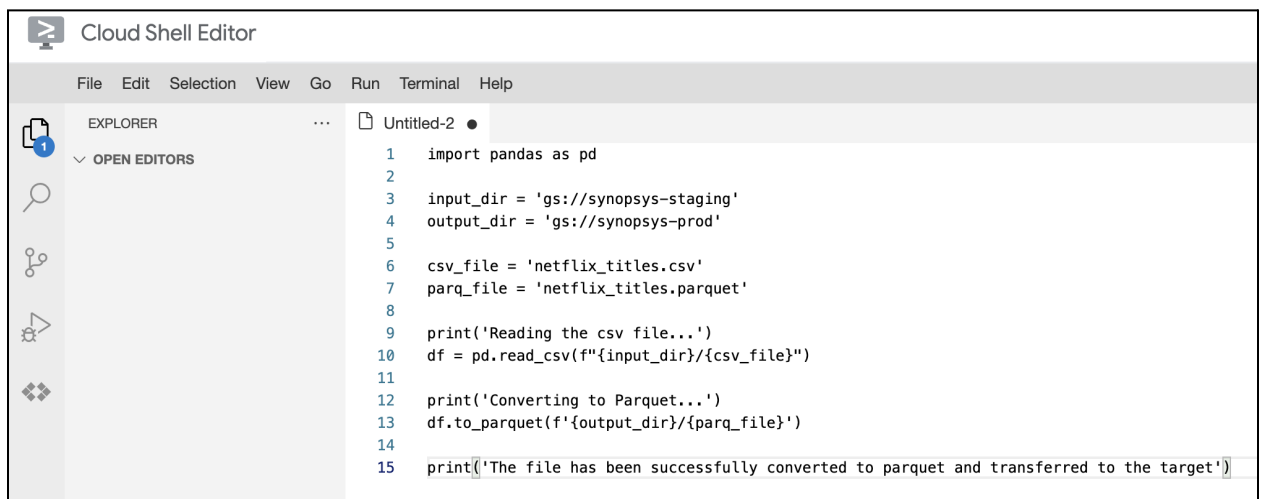
input_dir = 'gs://synopsys-staging'
output_dir = 'gs://synopsys-prod'

csv_file = 'netflix_titles.csv'
parq_file = 'netflix_titles.parquet'

print('Reading the csv file...')
df = pd.read_csv(f'{input_dir}/{csv_file}')

print('Converting to Parquet...')
df.to_parquet(f'{output_dir}/{parq_file}')

print('The file has been successfully converted to parquet and transferred to the target')
```



- Change the values for `input_dir` and `output_dir` to the staging and production buckets if necessary.
- Change the values for `csv_file` and `parq_file` accordingly with the proper extensions if necessary.
- Run the script using the following command and you'll find the new parquet file on the prod bucket

```
python csv_parquet.py
```

This is how your output should look like.

The screenshot shows the Cloud Shell Editor interface. The Explorer pane on the left shows the file structure. The main editor pane displays the `csv_parquet.py` script. The script imports `pandas` as `pd`, sets `input_dir` to `'gs://synopsys-staging2'` and `output_dir` to `'gs://synopsys-prod2'`. It then reads a CSV file `netflix_titles.csv` from the input directory, converts it to Parquet format, and transfers the resulting `netflix_titles.parquet` file to the output directory. The script includes print statements for each step and a final confirmation message.

```
1 import pandas as pd
2
3 input_dir = 'gs://synopsys-staging2'
4 output_dir = 'gs://synopsys-prod2'
5
6 csv_file = 'netflix_titles.csv'
7 parq_file = 'netflix_titles.parquet'
8
9 print('Reading the csv file...')
10 df = pd.read_csv(f'{input_dir}/{csv_file}')
11
12 print('Converting to Parquet...')
13 df.to_parquet(f'{output_dir}/{parq_file}')
14
15 print('The file has been successfully converted to parquet and transferred to the target')
16
```

The terminal output shows the successful execution of the script:

```
@cloudshell:~$ python csv_parquet.py
Reading the csv file...
Converting to Parquet...
The file has been successfully converted to parquet and transferred to the target
```

Once the file is successfully converted and transferred, you'll find it in the 'synopsys-prod' bucket.

The screenshot shows the Google Cloud Storage console for the 'synopsys-prod' bucket. The bucket is located in 'asia-south1 (Mumbai)' with a 'Standard' storage class, 'Not public' access, and 'None' protection. The 'OBJECTS' tab is selected, showing a list of objects. The object 'netflix_titles.parquet' is listed with a size of 1.9 MB, type 'application/octet-stream', and was created on Nov 15, 2023. It is stored in the 'Standard' storage class and was last modified on Nov 15, 2023. The public access is 'Not public' and it is encrypted with a 'Google-managed key'.

Location	Storage class	Public access	Protection
asia-south1 (Mumbai)	Standard	Not public	None

Objects	Configuration	Permissions	Protection	Lifecycle	Observability
Buckets > synopsys-prod Upload Files Upload Folder Create Folder Transfer Data Manage Holds Download Delete					

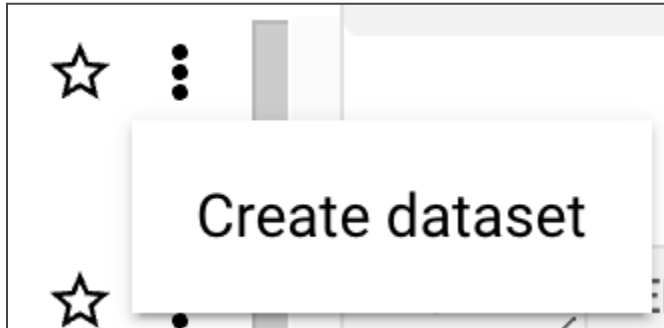
Name	Size	Type	Created	Storage class	Last modified	Public access	Version history	Encryption
netflix_titles.parquet	1.9 MB	application/octet-stream	Nov 15, 2023	Standard	Nov 15, 2023	Not public	—	Google-managed key

Step 2: GCS to BQ

Now that we have the Parquet file ready on a GCS bucket, we will be transferring it to BigQuery using the BigQuery Data Transfer Service.

First we need to **create a dataset in BigQuery**.

- A. Open the BigQuery page in the Google Cloud console.
- B. Go to the [BigQuery page](#)
- C. Click on SQL Workshop.
- D. In the Explorer panel, select the project where you want to create the dataset.
- E. Expand the Actions (three vertical dots) option next to your project name and click Create dataset.



- F. On the Create dataset page:
 - a. For Dataset ID, enter a unique dataset name, in this case we use `synopsys_demo`.
 - b. For Data location, choose a geographic location for the dataset. **This location should be the same as the one you used for creating the GCS buckets.**

Leave the other options as Default and click CREATE DATASET.

Create dataset

Project ID
searce-playground-v1

CHANGE

Dataset ID *
synopsys_demo

Letters, numbers, and underscores allowed

Data location
asia-south1 (Mumbai)

▼ ?

Default table expiration

☐ Enable table expiration ?

Default maximum table age

Days

Advanced options

▼

CREATE DATASET

CANCEL

Now we need to create an empty table in that dataset.

Run the following query in the Query **Editor** to create a new table 'netflix_title' in our Dataset. Make sure you replace PROJECT_ID with the ID of your GCP project.

Copy and paste the below query in SQL Workshop editor.

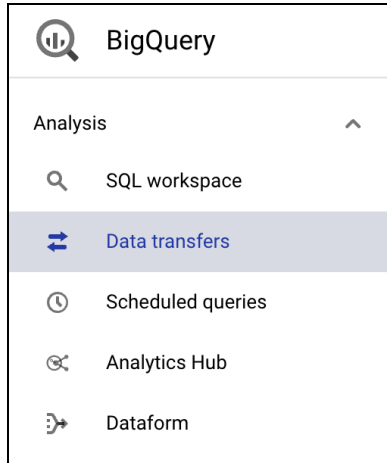
```
CREATE TABLE `PROJECT_ID.synopsys_demo.netflix_titles`  
(  
  show_id STRING,  
  type STRING,  
  title STRING,  
  director STRING,  
  `cast` STRING,  
  country STRING,  
  date_added STRING,  
  release_year INTEGER,  
  rating STRING,  
  duration STRING,  
  listed_in STRING,  
  description STRING  
);
```

Note: We are using backticks (") for `cast` because it is a reserved keyword in BigQuery

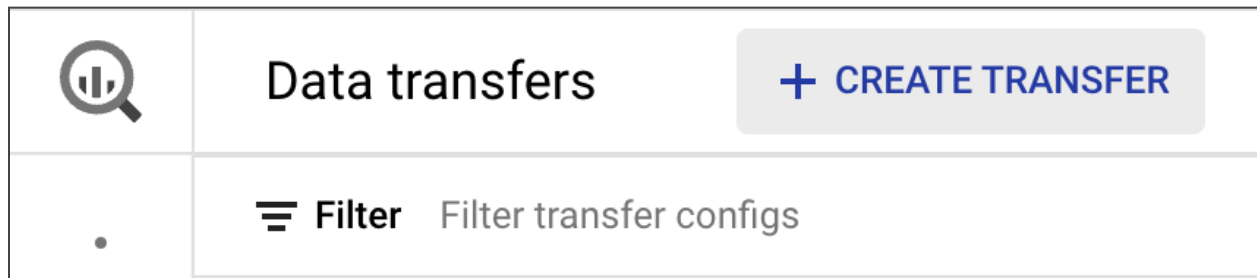
There are two ways to run the BigQuery Transfer. One is via using the console, and the other is using the CLI.

Using the Console - BQ Transfer

Open the Data Transfer Dashboard from BigQuery



Select “Create Transfer”.




Enter the configuration as follows:

Source: **Google Cloud Storage** as our source file is in a GCS bucket.

Display name: **synopsys-bq-transfer**



Repeats: **On-demand**

 Create transfer

Source type

Choose a data source from the list below

Source *
Google Cloud Storage ▼

 This is the Google Cloud Storage configuration. [Learn more](#) 


Transfer config name


Display name *
synopsys-bq-transfer

Schedule options

Repeats *
On-demand ▼

☒ Start now ☐ Start at set time

Start date and run time
11/15/22, 4:58 PM IST 

 Automatic scheduling is disabled. Transfer runs will be started on-demand via API or transfer details web page.

If you want the transfer to be repeated according to a fixed Schedule, you can select one of the other options.

Schedule options

Repeats *

Daily

Weekly

Monthly

Custom

On-demand

In this case, we are only going to make a one time transfer so we can go with **On-demand**. This way we can run the transfer whenever we want.

Destination settings

Select the destination for the transfer data

Dataset *

synopsys_demo

Data source details

Destination table

netflix_titles



Cloud Storage URI

 synopsys-prod/netflix_titles.parquet

BROWSE



Write preference

APPEND



☐ Delete source files after transfer 

File format *

PARQUET



Dataset: **synopsys_demo** name of the dataset we created.

Destination table: **netflix_titles** name of the table we created.

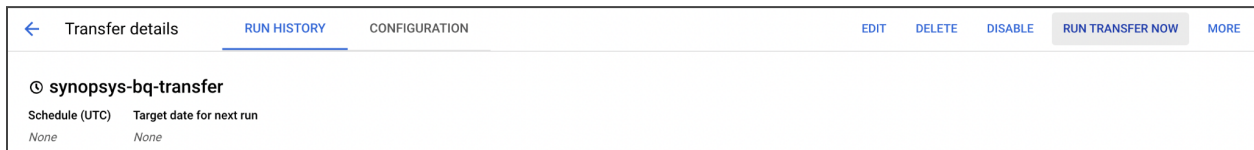
Cloud Storage URI: **synopsys-prod/netflix_titles.parquet** GCS URI of the source file. You can use the BROWSE option to select the file.

Write preference: **APPEND** as we will be adding the data to an existing table.
File format: **PARQUET** as our source file is in Parquet format.

Leave the other options as Default and click the SAVE button.

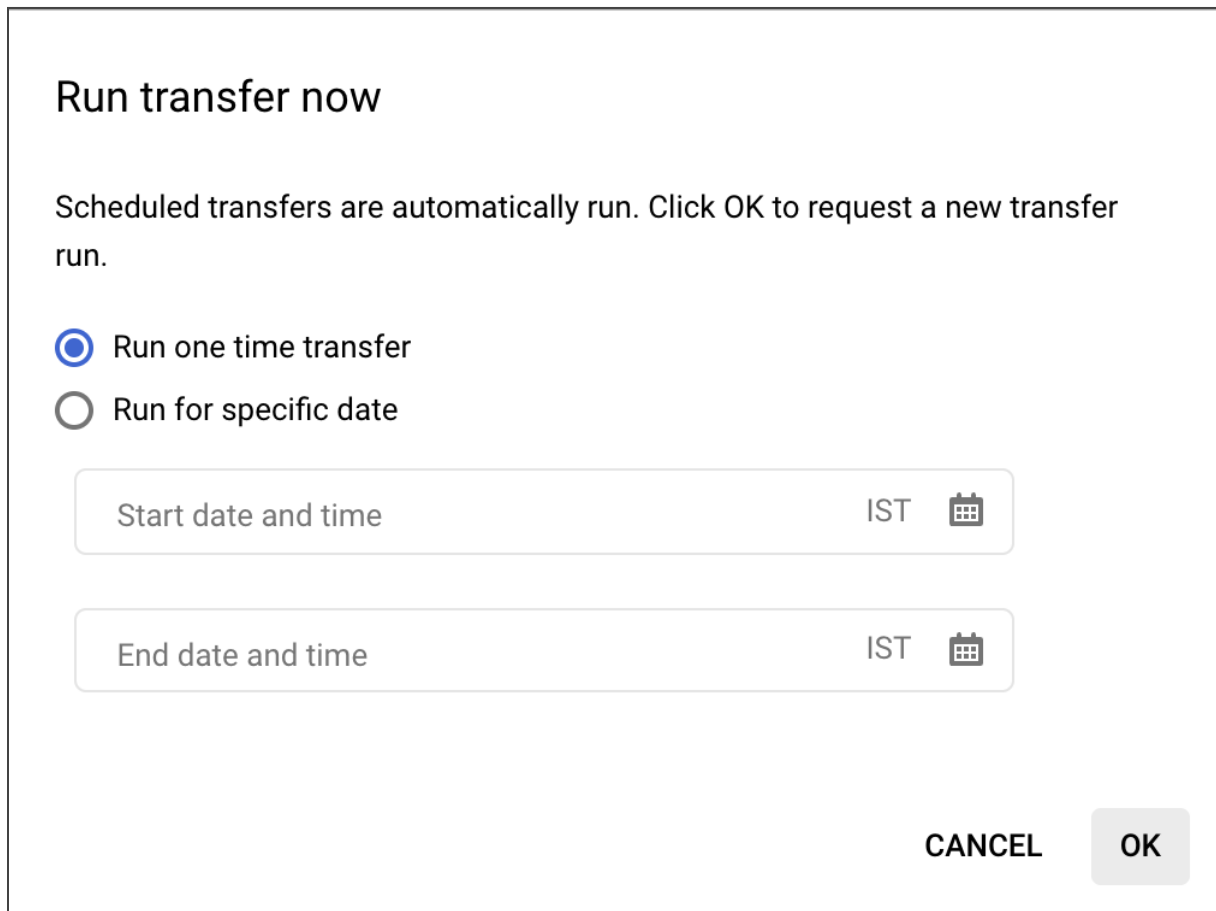
Run the Transfer job

Click the RUN TRANSFER NOW button on the top right corner to begin the transfer process.



The screenshot shows a configuration page for a transfer job named 'synopsys-bq-transfer'. The page has tabs for 'Transfer details', 'RUN HISTORY', and 'CONFIGURATION'. The 'Transfer details' tab is active. At the top right, there are buttons for 'EDIT', 'DELETE', 'DISABLE', 'RUN TRANSFER NOW', and 'MORE'. Below the job name, there are two fields: 'Schedule (UTC)' and 'Target date for next run', both with a value of 'None'.

We are only going to run it for one time, so select the first option and click OK



The screenshot shows a dialog box titled 'Run transfer now'. It contains the text: 'Scheduled transfers are automatically run. Click OK to request a new transfer run.' Below this, there are two radio button options: 'Run one time transfer' (which is selected) and 'Run for specific date'. Under the 'Run for specific date' option, there are two input fields: 'Start date and time' and 'End date and time', each followed by 'IST' and a calendar icon. At the bottom right, there are two buttons: 'CANCEL' and 'OK'.

Once the transfer is complete, you can find the resulting table on BQ

synopsys-bq-transfer

Schedule (UTC)

Target date for next run

None

None

Filter

Filter transfer runs

		Run date ? ↓	Schedule time UTC+5:30 ?	Summary
		November 15, 2022	November 15, 2022 at 5:25:06 PM UTC+5:30	The transfer run has completed successfully.

To view the BQ Table :

Go to the Explorer Panel, first expand the project and then expand the dataset you created, and then select the table

Click on Preview, this is how the BQ table should look like

netflix_titles

<

From the DETAILS tab, you can see that we have successfully added all the 8,807 rows and the table size is 3.29 MB.

Storage info

Number of rows	8,807
Total logical bytes	3.29 MB
Active logical bytes	3.29 MB
Long term logical bytes	0 B
Total physical bytes	0 B
Active physical bytes	0 B
Long term physical bytes	0 B
Time travel physical bytes	0 B

Method 2: Using CLI

The transfer can also be done using the 'bq load' CLI command

```
bq load \  
--source_format=PARQUET \  
synopsys_demo.netflix_titles \  
"gs://<bucket-name>/netflix_titles.parquet"
```

Note : Replace `<bucket-name>` with your prod bucket name

The parameters used are:

1. The file format of the source file, in this case it is **PARQUET**
2. The destination table in the form 'DATASET.TABLE_NAME', in this case, its **synopsys_demo.netflix_titles**
3. The URI of the source file, in this case its **gs://synopsys-prod/netflix_titles.parquet**

Once the job has successfully run, you'll find the following output.

```
🍏 > ~/DE/Pro/synopsys bq load \  
--source_format=PARQUET \  
synopsys_demo.netflix_titles \  
"gs://synopsys-prod/netflix_titles.parquet"  
Waiting on bqjob_r362ee906afbe0c2_000001847bd9c658_1 ... (1s) Current status: DONE
```


The disadvantage of using this method is that you can't schedule the transfers to run at particular times.

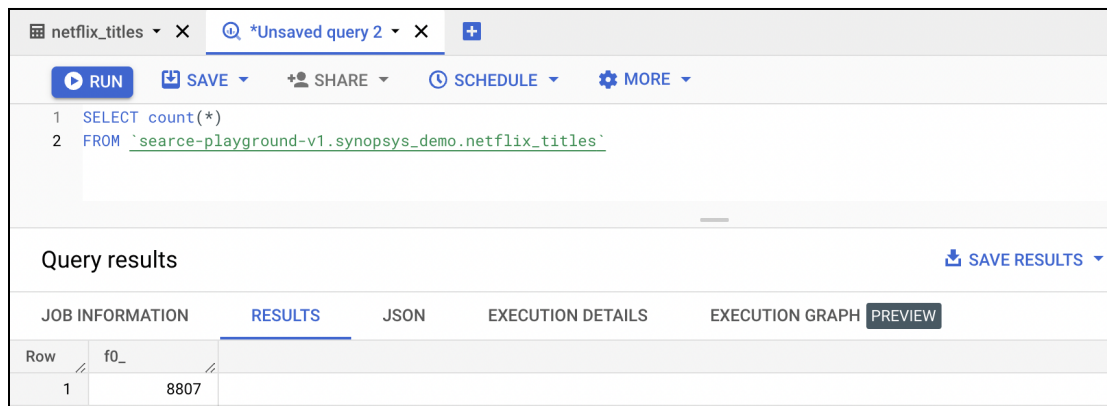
Step 3 : BQ sample Queries

- a. Total number of records

Query

```
SELECT count(*)
FROM `<project_id>.synopsys_demo.netflix_titles`
```

Results



The screenshot shows the Google BigQuery web interface. At the top, there are tabs for 'netflix_titles' and '*Unsaved query 2'. Below the tabs are buttons for 'RUN', 'SAVE', 'SHARE', 'SCHEDULE', and 'MORE'. The query editor contains the following SQL code:

```
1 SELECT count(*)
2 FROM `searce-playground-v1.synopsys_demo.netflix_titles`
```

Below the query editor, the 'Query results' section is visible. It includes a 'SAVE RESULTS' button and tabs for 'JOB INFORMATION', 'RESULTS', 'JSON', 'EXECUTION DETAILS', 'EXECUTION GRAPH', and 'PREVIEW'. The 'RESULTS' tab is selected, showing a table with one row and one column:

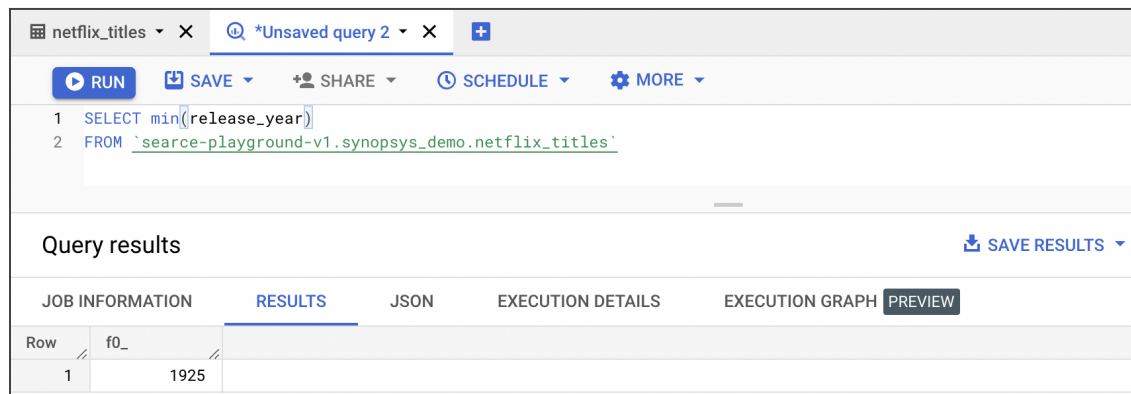
Row	f0_
1	8807

- b. Find the year in which the oldest title was released

Query

```
SELECT min(release_year)
FROM `<project_id>.synopsys_demo.netflix_titles`
```

Results



The screenshot shows the Google BigQuery web interface. At the top, there are tabs for 'netflix_titles' and '*Unsaved query 2'. Below the tabs are buttons for 'RUN', 'SAVE', 'SHARE', 'SCHEDULE', and 'MORE'. The query editor contains the following SQL code:

```
1 SELECT min(release_year)
2 FROM `searce-playground-v1.synopsys_demo.netflix_titles`
```

Below the query editor, the 'Query results' section is visible. It includes a 'SAVE RESULTS' button and tabs for 'JOB INFORMATION', 'RESULTS', 'JSON', 'EXECUTION DETAILS', 'EXECUTION GRAPH', and 'PREVIEW'. The 'RESULTS' tab is selected, showing a table with one row and one column:

Row	f0_
1	1925

c. Number of titles from each director ordered from greatest to the least
Query

```
SELECT director, count(*) AS titles
FROM `<project_id>.synopsys_demo.netflix_titles`
GROUP BY director
ORDER BY titles DESC
```

Results

netflix_titles

*Unsaved query 2

RUN

SAVE

SHARE

SCHEDULE

MORE

```
1 SELECT director, count(*) AS titles
2 FROM `searce-playground-v1.synopsys_demo.netflix_titles`
3 GROUP BY director
4 ORDER BY titles DESC
```

Query results

JOB INFORMATION

RESULTS

JSON

EXECUTION DETAILS

EXECUTION GRAPH

PREVIEW

Row	director	titles
2	Rajiv Chilaka	19
3	Raúl Campos, Jan Suter	18
4	Suhas Kadav	16
5	Marcus Raboy	16
6	Jay Karas	14
7	Cathy Garcia-Molina	13
8	Martin Scorsese	12
9	Jay Chapman	12
10	Youssef Chahine	12

d. Title count for each category in listed_in ordered from highest to lowest
Query

```
SELECT listed_in, count(*) AS titles
FROM `<project_id>.synopsys_demo.netflix_titles`
GROUP BY listed_in
ORDER BY titles DESC
```

Results

netflix_titles

*Unsaved query 2

RUN

SAVE

SHARE

SCHEDULE

MORE

```
1 SELECT listed_in, count(*) AS titles
2 FROM `search-playground-v1.synopsys_demo.netflix_titles`
3 GROUP BY listed_in
4 ORDER BY titles DESC
5
```

Query results

JOB INFORMATION

RESULTS

JSON

EXECUTION DETAILS

EXECUTION GRAPH

PREVIEW

Row	listed_in	titles
1	Dramas, International Movies	362
2	Documentaries	359
3	Stand-Up Comedy	334
4	Comedies, Dramas, Internation...	274
5	Dramas, Independent Movies, I...	252
6	Kids' TV	220
7	Children & Family Movies	215
8	Children & Family Movies, Com...	201
9	Documentaries, International M...	186
10	Dramas, International Movies, ...	180

Use Cases

The use cases most suited for BigQuery are the ones where the users need to perform interactive ad-hoc queries of read-only datasets.

Typically, BigQuery is used at the end of a Big Data ETL pipeline on top of processed data or in scenarios where complex analytical queries to a relational database take several seconds.

As it has a built-in cache, BigQuery works really well in cases where the data does not change often.

In addition, cases where datasets are fairly small don't really benefit from BigQuery, with a simple query taking up to a few seconds. As such, it shouldn't be used as a regular OLTP (Online Transaction Processing) database.

BigQuery is meant and suited for OLAP (Online Analytics Processing) Analytics workloads.

Why BigQuery Transfer Service?

An alternative method to BQ Transfer Service for transferring data from GCS to BigQuery is using Cloud Dataflow.

Both methods have their advantages and disadvantages.

Method 1: Using BigQuery Data Transfer Service

Pros:

- The entire process is done on the console and no coding is involved.
- The prerequisites and IAM permissions required are minimal.
- The method is simple and straightforward. Instructions are easy to follow and less chances of getting unexpected errors.

Cons:

- Larger files (>100MB) need to be divided/sharded into multiple smaller ones.
- We cannot perform any transformations on the data before it is loaded into BQ.

Method 2: Using a Custom Dataflow Pipeline

Pros:

- Can handle large files with the appropriate configuration.
- Much more customisable.
- We can perform custom transformations on the data before it is loaded into BQ.

Cons:

- Requires many IAM permissions along with a Service Account.
- More prerequisites are required both in terms of knowledge and software packages
- Relatively more advanced and not beginner friendly, especially to someone who has no familiarity with Dataflow or Apache Beam.
- As it is a complicated process, there is a lot of room for unexpected errors and blockers to occur, which can be hard to navigate for someone who has no expertise in this.

For this tutorial we went with Method 1 because for our use case, the file size is relatively small and we did not have to make any transformations to the data before loading it to BigQuery

Best Practices

1. Use GCS as staging area before uploading to BigQuery

In many use cases including this tutorial, it is recommended to use GCS as a staging area before transferring it to BigQuery. Here are the pros and cons of doing so.

Pros

- Loading data is incredibly fast
- Possible to remove BQ table when not used and import it when needed (BQ table can be much bigger than even compressed data in Cloud Storage)
- You save your local storage space
- Less likely to fail during table creation (from local storage there could be networking issues, computer issues etc.)

Cons

- You pay some additional cost for storage (in the case you do not plan to access the data often e.g. once per month or once a year - you can reduce costs by using the nearline or coldline storage types in GCS)

2. Nested and Repeated Data

This is one of the most important Google BigQuery ETL best practices. Google BigQuery performs best when the data is denormalized. Instead of keeping relations, denormalize the data and take advantage of nested and repeated fields.

Nested and repeated fields are supported in Avro, Parquet, ORC, JSON (newline delimited) formats. STRUCT is the type that can be used to represent an object which can be nested and ARRAY is the type to be used for the repeated value.

3. Data Compression

The next vital Google BigQuery ETL best practice is on Data Compression. Most of the time the data will be compressed before transfer. You should consider the below points while compressing data.

- The binary Avro is the most efficient format for loading compressed data.
- Parquet and ORC format are also good as they can be loaded in parallel.
- For CSV and JSON, Google BigQuery can load uncompressed files significantly faster than compressed files because uncompressed files can be read in parallel.

Always try to convert large CSV or JSON files into file types such as Avro or Parquet before transferring to BigQuery. This will save significant time, resources, and costs.

4. Access Control and Data Encryption

Data stored in Google BigQuery is encrypted by default and keys are managed by GCP. Alternatively customers can manage keys using the Google KMS service.

To grant access to resources, BigQuery uses IAM (Identity and Access Management) to the dataset level. Tables and views are child resources of datasets and

inherit permission from the dataset. There are predefined roles like `bigquery.dataViewer` and `bigquery.dataEditor` or the user can create custom roles.

Make sure all users are assigned the minimum required roles/permissions at a granular level.

5. Bulk Updates

When you are making UPDATES to your BQ table, make them in bulk instead of multiple separate queries.

BigQuery has quotas and limits for DML statements which are getting increased over time. As of now the limit of combined INSERT, UPDATE, DELETE and MERGE statements per day per table is 1,000.

Note that this is not the number of rows. This is the number of the statements and as you know, one single DML statement can affect millions of rows.

Now within this limit, you can run updates or merge statements affecting any number of rows. It will not affect any query performance, unlike many other analytical solutions.